

Basic Statistics for Data Science with Python

Nur Andi Setiabudi

July 9, 2021



Food Tasting

- Stir it well
- Pick only small portion of it
- Taste
- Conclude

That was
Statistics

MATCH STATISTICS

QUARTER-FINALS - FIRST LEG



LIVERPOOL



FULL TIME

2 - 0

PORTO



58%

3

14

1

4

4

107.45 km

539 (82%)

9

0 / 0

BALL POSSESSION

ATTEMPTS ON TARGET

TOTAL ATTEMPTS

SAVES

CORNERS

OFFSIDES

DISTANCE COVERED

PASSES COMPLETED

FOULS COMMITTED

YELLOW/RED CARDS

42%

3

9

3

5

2

109.26 km

253 (68%)

7

2 / 0

01:00 WIB



Cerah Berawan

22°C

90 %

10 km/jam

Timur

04:00 WIB



Cerah Berawan

20°C

95 %

10 km/jam

Timur

07:00 WIB



Cerah Berawan

24°C

80 %

10 km/jam

Timur

10:00 WIB



Cerah

28°C

60 %

10 km/jam

Timur

13:00 WIB



Hujan Sedang

27°C

60 %

30 km/jam

Timur

16:00 WIB



Hujan Sedang

28°C

60 %

20 km/jam

Timur

19:00 WIB



Hujan Ringan

24°C

80 %

10 km/jam

Timur

22:00 WIB



Berawan



24°C

85 %

10 km/jam

Timur

Frequently Bought Together



Total price: \$83.09

[Add both to Cart](#)

[Add both to List](#)

[This item: Structure and Interpretation of Computer Programs - 2nd Edition \(MIT Electrical Engineering and... by Harold Abelson Paperback \\$50.50](#)

[The Pragmatic Programmer: From Journeyman to Master by Andrew Hunt Paperback \\$32.59](#)

Customers Who Bought This Item Also Bought



The Little Schemer - 4th Edition
by Daniel P. Friedman
★★★★☆ 64
Paperback \$36.00 ✓Prime



Instructor's Manual Va Structure and Interpretation of Computer Programs...
by Gerald Jay Sussman
★★★★☆ 5
Paperback \$28.70 ✓Prime



The Pragmatic Programmer: From Journeyman to Master
by Andrew Hunt
★★★★☆ 328
Paperback \$32.59 ✓Prime





Introduction to Algorithms, 3rd Edition (MIT Press)
by Thomas H. Cormen
★★★★☆ 313
#1 Best Seller in Computer Algorithms
Hardcover \$66.32 ✓Prime



An Introduction to Functional Programming Through Lambda Calculus...
by Greg Michaelson
★★★★☆ 23
Paperback \$20.70 ✓Prime



Purely Functional Data Structures
by Chris Okasaki
★★★★☆ 19
Paperback \$40.74 ✓Prime






*But, statistics (and also data science, of course) **is not a magic***



Definition



What is statistics?


- The practice or **science** of **collecting** and **analyzing** numerical **data** in large quantities, especially for the purpose of **inferring proportions** in a whole from those in a **representative sample** – Oxford
- The **science** of learning from **data**, and of **measuring**, **controlling** and **communicating uncertainty** – American Statistical Association

Basic steps in statistics

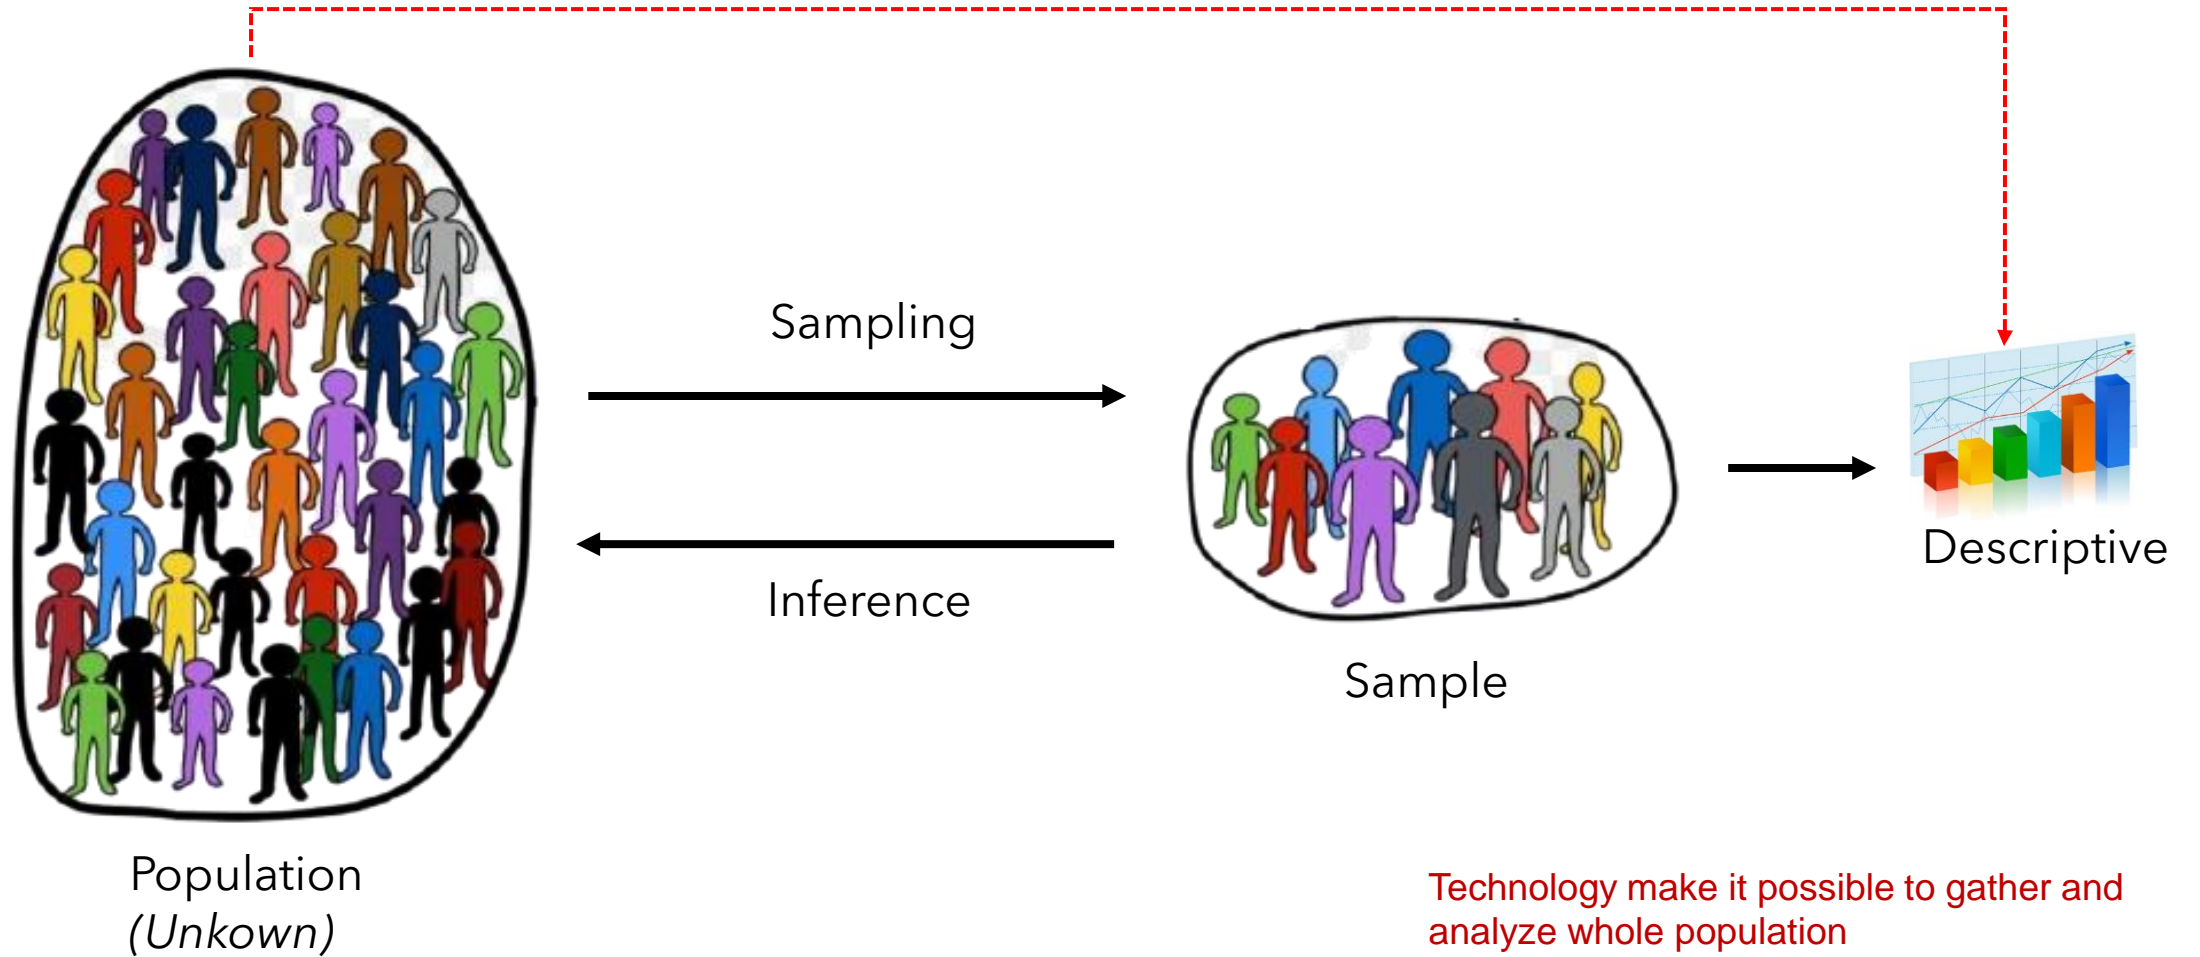
Studying a problem through the use of statistical data analysis usually involves four basic steps.

- Defining the problem
- Collecting/preparing the data
- Analyzing the data
- Reporting the results

Population and Sample



Statistics



Population and sample

- Data consists of information coming from observations, counts, measurements, or responses
- A population is the entire group that you want to draw conclusions about.
- A sample is the specific group that you will collect data from. The size of the sample is always less than the total size of the population.

Sample
Statistics



Population
Parameter


- A parameter is a numerical description of a *population* characteristic.
- A statistic is a numerical description of a *sample* characteristic.

Characteristics of good sample


- **Representative** of population: should be an accurate representative of the universe from which it is taken
- **Random selection**: should be selected at random. This means that any item in the group has a full and equal chance of being selected and included in the sample. This makes the selected sample truly representative
- **Sampling error** can be quantified. **Non sampling error** is corrected for as much as possible
- **Economical**: should be achieved with minimum cost and effort
- **Practical**: should be capable of being understood and followed in the fieldwork

Sampling method


Simple random sample




Systematic sample




Convenience sample




Voluntary response sample




Stratified sample




Cluster sample



Purposive sample



Snowball sample



Error and bias

Potential source of error

in estimating population parameter using sample

Sampling
error

Sample is not
the whole
population


Non-sampling error

Behavioral
effect

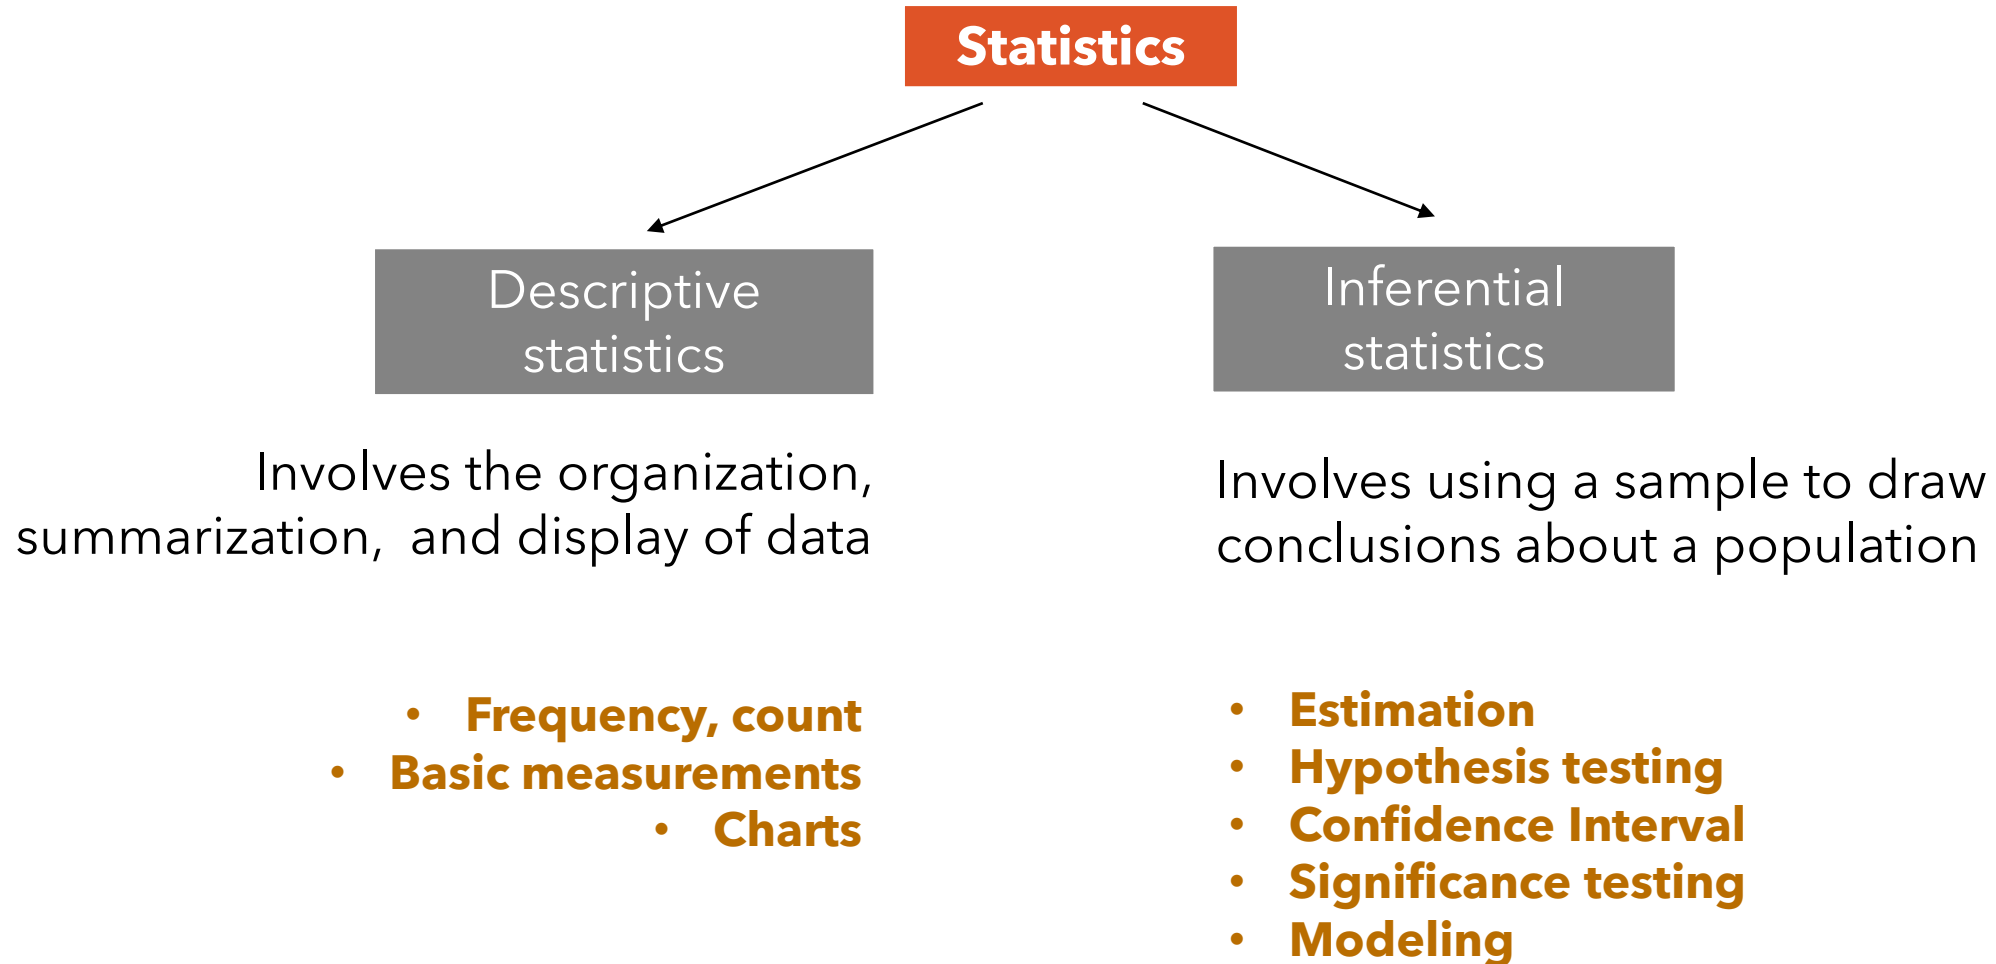
Questionnaire
and tools error

Poor sampling
method


Error vs sample size



Two Branches of Statistics



I Descriptive Statistics



Measurement level

Qualitative		Quantitative	
Nominal	Ordinal	Interval	Ratio
Named	Named	Named	Named
Color Gender	Natural order	Natural order	Natural order
	Education Satisfaction level	Equal interval between variables	Equal interval between variables
		Temperature Exam score	Has a "true zero" value, thus ratio between values can be calculated
			Age Monthly income

Descriptive statistics

- Descriptive statistics is the term given to the analysis of data that helps describe, show or summarize data in a meaningful way such that, for example, patterns might emerge from the data.
- Descriptive statistics **do not**, however, allow us to make conclusions beyond the data we have analyzed or reach conclusions regarding any hypotheses we might have made.
- They are simply a way to describe our data.
- **Measures of central tendency:** these are ways of describing the central position of a frequency distribution for a group of data
- **Measures of spread (dispersion):** these are ways of summarizing a group of data by describing how spread out the scores are

Measures of central tendency

ways of describing the central position of a frequency distribution for a group of data


Measure of Central Tendency	Appropriate to choose when ...	Should not be used when...
Mean the balance point of a data distribution	No situation precludes it	<ul style="list-style-type: none">• Extreme scores• Skewed distribution• Ordinal scale• Nominal scale
Median the midpoint of a data distribution	<ul style="list-style-type: none">• Extreme scores• Skewed distribution• Ordinal scale	<ul style="list-style-type: none">• Nominal scale
Mode score or category that has the greatest frequency	<ul style="list-style-type: none">• Nominal scales• Discrete variables• Describing shape	<ul style="list-style-type: none">• Interval or ratio data, except to accompany mean or median

Measures of dispersion



Summarizing a group of data by describing how spread out the scores are

- **Range:** The interval between the highest and lowest measures
- **Percentile:** The value below / above which a particular percentage of values fall
- The **standard deviation** is a statistic that measures the dispersion of a dataset relative to its mean and is calculated as the square root of the variance

Data Distribution




Skewness



Box Plot

Presenting data: Table and chart

Age	Tally	Frequency
1-10		5
11-20	III	8
21-30		14
31-40		18
41-50		20
51-60		13
61-70	I	6



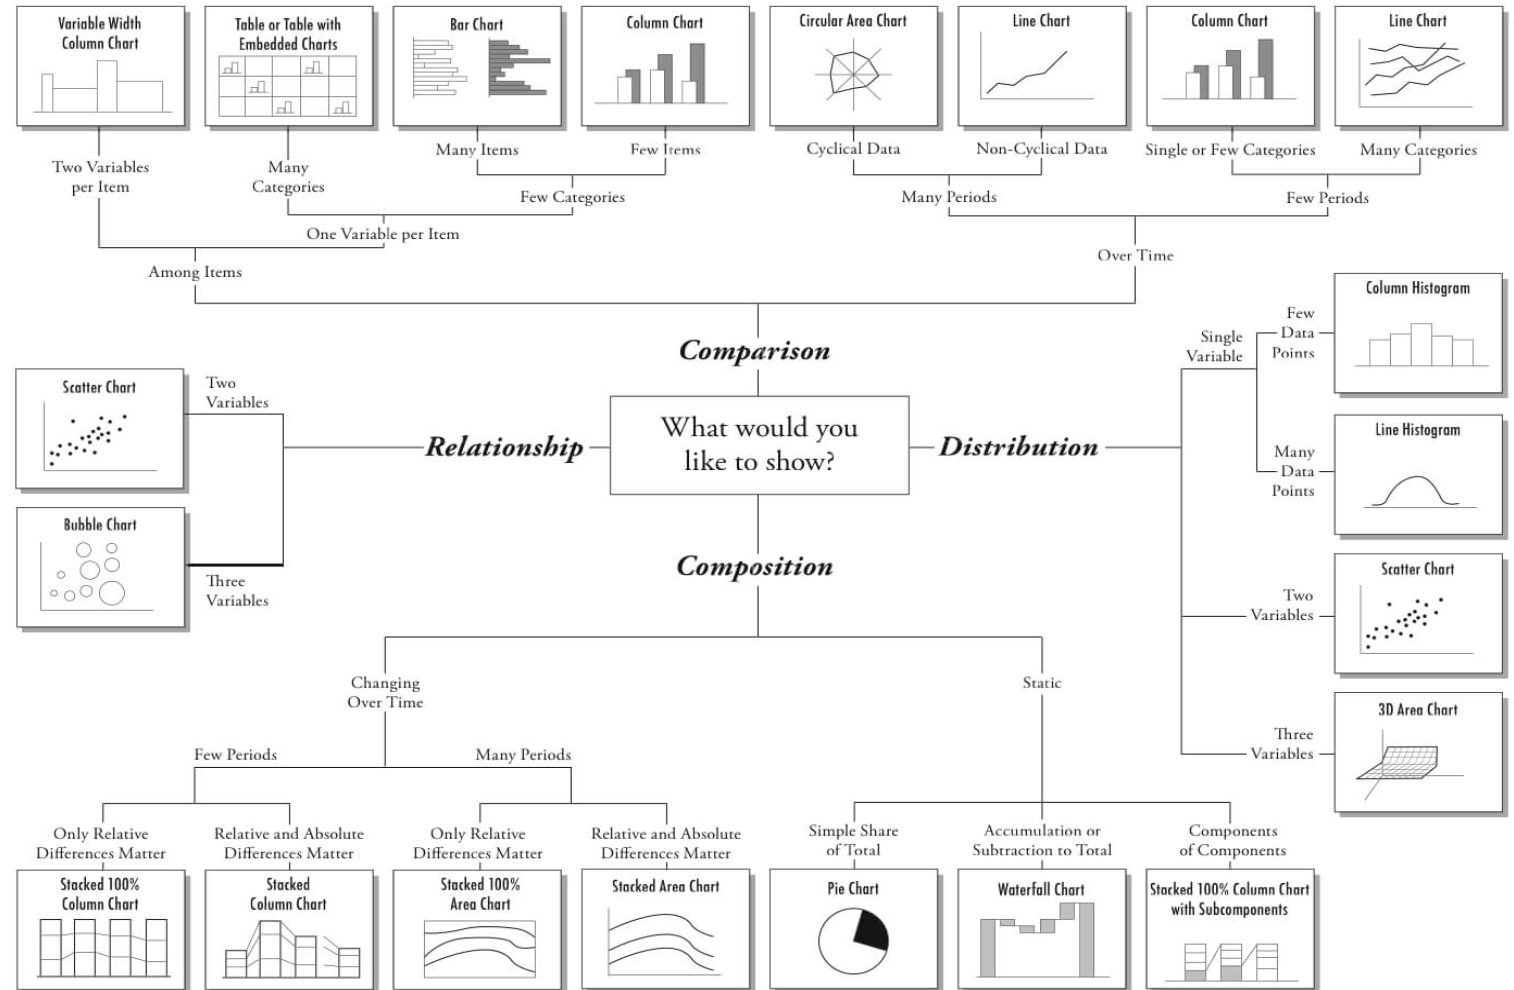
There is no space between bars.

Because the intervals are equal, all of the bars have the same width.

Presenting data: chart


Chart Suggestions—A Thought-Starter

www.ExtremePresentation.com
© 2009 A. Abela — a.v.abela@gmail.com





Probability Distribution



Probability distribution terminology

- Probability is the measure of the likelihood that an event will occur in a random experiment. It is quantified as a number between 0 and 1
- A random experiment is a physical situation whose outcome cannot be predicted until it is observed
- A sample space, is a set of all possible outcomes of a random experiment


Probability distribution terminology

- A sample will form a **distribution**
- The distribution provides a **parameterized mathematical function** that can be used to calculate the probability for any individual observation from the sample space
- Data type:
 - Discrete: take only specified values
 - Continuous: take any value within a given range

Common data distribution

- Many data conform to well-known and well-understood mathematical functions
 - **Bernoulli**: two possible values, eg. 0/1, success/fail, etc
 - **Uniform**: finite number of values are equally likely to be observed, eg. rolling dice
 - **Binomial**: n times random experiment of Bernoulli
 - **Normal**: represents the behavior of most of the situations in the universe
 - Chi-square, Poisson, Gamma, Exponential, etc

Normal (**z**) distribution




The Normal Distribution has:

- mean = median = mode
- symmetry about the center
- 50% of values less than the mean
- and 50% greater than the mean

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2}$$


Student's t distribution

- Like a standard normal distribution (or z -distribution), the t -distribution has a mean of zero.
- The normal distribution assumes that the population standard deviation is known. The t -distribution does not make this assumption.
- The t -distribution is defined by the *degrees of freedom*. These are related to the sample size.
- The t -distribution is most useful for small sample sizes, when the population standard deviation is not known, or both.
- As the sample size increases, the t -distribution becomes more similar to a normal distribution.




Standardized Normal distribution

**Normal
distribution**



$$Z = \frac{x - \mu}{\sigma}$$


**Standardized
Normal distribution**



Area under curve = 1
(probability of all-possible events)

Calculate probability: **example**


Problem: The weights of adult-males are known to be normally distributed with a mean of 70 kgs and a standard deviation of 13 kgs. Find the percentage of adult-males with weights less than 80 kgs



For $X = 80 \rightarrow Z = (80 - 70)/13 = 0.769$

So, $P(X < 80) = P(Z < 0.769)$

Calculate probability: **example**



$$P(\mathbf{X} < 80) = P(\mathbf{Z} < 0.769) = 0.7794$$


The percentage of adult-males with weights less than 80 kgs is around 78%

Standard Normal Probabilities

Table entry for z is the area under the standard normal curve to the left of z .

z	.00	.01	.02	.03	.04	.05	.06	.07	.08	.09
0.0	.5000	.5040	.5080	.5120	.5160	.5199	.5239	.5279	.5319	.5359
0.1	.5398	.5438	.5478	.5517	.5557	.5596	.5636	.5675	.5714	.5753
0.2	.5793	.5832	.5871	.5910	.5948	.5987	.6026	.6064	.6103	.6141
0.3	.6179	.6217	.6255	.6293	.6331	.6368	.6406	.6443	.6480	.6517
0.4	.6554	.6591	.6628	.6664	.6700	.6736	.6772	.6808	.6844	.6879
0.5	.6915	.6950	.6985	.7019	.7054	.7088	.7123	.7157	.7190	.7224
0.6	.7257	.7291	.7324	.7357	.7389	.7422	.7454	.7486	.7517	.7549
0.7	.7580	.7611	.7642	.7673	.7704	.7734	.7764	.7794	.7823	.7852
0.8	.7881	.7910	.7939	.7967	.7995	.8023	.8051	.8078	.8106	.8133
0.9	.8159	.8186	.8212	.8238	.8264	.8289	.8315	.8340	.8365	.8389
1.0	.8413	.8438	.8461	.8485	.8508	.8531	.8554	.8577	.8599	.8621
1.1	.8643	.8665	.8686	.8708	.8729	.8749	.8770	.8790	.8810	.8830
1.2	.8849	.8869	.8888	.8907	.8925	.8944	.8962	.8980	.8997	.9015
1.3	.9032	.9049	.9066	.9082	.9099	.9115	.9131	.9147	.9162	.9177

I Inferential statistics



Inferential statistics



Eight **out of 10** owners said their **cat** prefers it



80% ?

How confident can we be about such statistic?

8 out of 10?


80 out of 100?

800 out of 1000?

80000 out of 100000?


Inferential statistics

- Process to draw conclusions about some unknown aspect of a population based on a random sample from that population



Confidence interval

- A confidence interval calculates the probability that a population parameter will fall between two set values.
- Confidence intervals measure the degree of uncertainty or certainty in a sampling method.
- Most often, confidence intervals reflect confidence levels of 95% or 99%.
- If we say that we are 95% confident that the unknown population mean is contained in this interval, then we are really saying that we found the interval using a method that is successful in giving correct results 95% of the time




Interpreting confidence interval

"In a sample of 659 parents with toddlers, about 85%, stated they use a car seat for all travel with their toddler. From these results, a 95% confidence interval was provided, going from about 82.3% up to 87.7%."

- we are 95% certain that the population proportion who use a car seat for all travel with their toddler will fall between 82.3% and 87.7%.
- if we take a different sample from these 659 people, 95% of the time, the percentage of the population who use a car seat in all travel with their toddlers will be in between 82.3% and 87.7%.
- **95% confidence interval does not mean 95% probability**

Confidence level vs confidence interval

The greater the confidence level, the wider the confidence interval



Confidence interval formula

- Confidence interval for mean

$$\bar{X} - Z_{\alpha} \left(\frac{\sigma}{\sqrt{n}} \right) \leq \mu \leq \bar{X} + Z_{\alpha} \left(\frac{\sigma}{\sqrt{n}} \right)$$

(Population variance known)

$$\bar{x} - t_{v,\alpha} \left(\frac{s}{\sqrt{n}} \right) \leq \mu \leq \bar{x} + t_{v,\alpha} \left(\frac{s}{\sqrt{n}} \right)$$

(Population variance unknown)


- Confidence interval for proportion

$$p' - Z_{\alpha} \sqrt{\frac{p'q'}{n}} \leq p \leq p' + Z_{\alpha} \sqrt{\frac{p'q'}{n}}$$

Confidence interval: **example**


Problem: We measure the heights of 40 randomly chosen men, and get a mean height of 175cm. We have already known that standard deviation is 20cm. Calculate the 95% confidence interval for mean height!

$$Z_{\alpha} \left(\frac{\sigma}{\sqrt{n}} \right) = 1.96 * (20/\sqrt{40}) = 6.2$$



This says the true mean of ALL men height is likely to be between 168.8cm and 181.2cm.

I Hypothesis testing



Hypothesis testing

- Hypothesis testing is a statistical method that is used in making statistical decisions using sample
- Hypothesis Testing is basically an assumption that we make about the population parameter
- Hypothesis is made before running the test

Null Hypothesis (H_0)

Assumes that the observation is due to a chance factor

"Two website designs generate equal revenue"

VS

Alternative Hypothesis (H_1)

Shows that observations are the result of a real effect
This is hypothesis that we propose

"New website design generate more revenue than the existing one"

Hypothesis testing ...

Null Hypothesis (H_0)

"Two website designs generate equal revenue"

VS

Alternative Hypothesis (H_1)

"New website design generate more revenue than the existing one"

- Collect **sufficient** evidence to reject $H_0 \rightarrow$ accept H_1
- How sufficient is sufficient? \rightarrow Confidence level (90%, 95%, 99%, ect.)
- When reject H_0 , we conclude that H_1 is true at certain level of confidence
- We never accept H_0 . Instead, we say: No sufficient evidence to reject H_0

Two types of error

	Reject H0	Fail to Reject H0
Reality: H0 is True	Type I error (α)	Correct decision Probability ($1 - \alpha$)
Reality: H0 is False	Correct decision Power ($1 - \beta$)	Type II error (β)

Hypothesis testing: the court



Man
About 40s
180 cm tall
Black sedan



Hypothesis testing: the court



Courtroom
hypothesis

Innocent

Our job is to disapprove
this

If innocent, very
unlikely to find
this evidence

until proven

EVIDENCE

Guilty

Alternative hypothesis

So we can accept this

HOW UNLIKELY?

50%: chance he's innocent
and we found that evidence

20%

5%

1%

Significant level
(alpha)

Hypothesis testing: the court



Judge doesn't reject
his innocent

**CORRECT
DECISION**



Judge doesn't reject
his innocent

ERROR (Type II) :
 β



Judge reject his
innocent

ERROR (Type I) :
 α



Judge reject his
innocent

**CORRECT
DECISION**



- There are un-avoidable errors in statistics
- That's why statistics is also defined as **science of uncertainty**
- The role of statisticians is to **quantify** and **minimize** the error

Statistical test

Correlational: test an association or relationship between variables

Pearson Correlation	Tests for the strength of the association between two continuous variables
Spearman Correlation	Tests for the strength of the association between two ordinal variables

Comparison of Means: test the difference between the means of variables

T-Test	Tests for the difference between two group
ANOVA	Tests for the difference between group means (more than 2 groups)

Regression: test if change in one variable predicts change in another variable

Linear regression	Tests how change in the predictor variable predicts the level of change in the outcome variable
Logistics regression	Tests how change in the predictor variable predicts the level of change in the categorical outcome variable

Hypothesis testing steps

1. State the null hypothesis, H_0 and the alternative hypothesis, H_1
2. Choose an alpha α , our significance level
3. Select a **statistical test**, and calculate the observed **test statistic**
4. Find the **critical value** of the test statistic (and/or **p-value**)
5. Compare the observed test statistic with the critical value, (or compare the **p-value** with α), and decide to accept or reject H_0

Two and one-tailed hypothesis testing

$$H_0: \mu = 15.00$$

$$H_1: \mu \neq 15.00$$

Two-tailed

$$H_0: \mu \geq 15$$


$$H_1: \mu < 15$$

One-tailed



$$H_0: \mu \leq 15$$

$$H_1: \mu > 15$$

Alpha/significance level



Alpha, critical value, test statistic and p-value



Making conclusion

If test statistic is large enough to be in the critical region, we conclude that the difference is significant or that the treatment has a significant effect. **In this case we reject the null hypothesis.**

If the mean difference is relatively small, then the test statistic will have a low value. In this case, we conclude that the **evidence from the sample is not sufficient**, and the decision is **fail to reject the null hypothesis**

Hypothesis testing about mean

- **One sample:**

- Z test if population variance is known

$$z_h = \frac{\bar{x} - \mu_0}{\sigma / \sqrt{n}}$$

- Z test if population variance is unknown

$$t_h = \frac{\bar{x} - \mu_0}{s / \sqrt{n}}$$

Hypothesis testing about mean

- **Two samples**

- Both variances are unknown, but assumed equal)

$$t_h = \frac{(\bar{x}_1 - \bar{x}_2) - \delta_0}{s_{(\bar{x}_1 - \bar{x}_2)}} \quad s_{\bar{x}_1 - \bar{x}_2} = \sqrt{s_{gab}^2 \left(\frac{1}{n_1} + \frac{1}{n_2} \right)}$$
$$s_{gab}^2 = \frac{(n_1 - 1)s_1^2 + (n_2 - 1)s_2^2}{n_1 + n_2 - 2} \quad \text{dan } v = n_1 + n_2 - 2$$

Hypothesis testing about mean

- **Two paired samples**

$$t = \frac{m}{s / \sqrt{n}}$$

Since ***m*** is mean of pairwise difference between two-sample, we can simply apply standard one-sample T-Test

Hypothesis testing about proportion

- **One sample proportion**

$$z = \frac{\hat{p} - p_0}{\sqrt{\frac{p_0(1-p_0)}{n}}}$$

- **Comparison of two-proportions**

$$z = \frac{(\hat{p}_1 - \hat{p}_2)}{\sqrt{\hat{p}(1-\hat{p})\left(\frac{1}{n_1} + \frac{1}{n_2}\right)}}$$

Levene's test

- Inferential statistic used to assess the **equality of variances** for a variable calculated for two or more groups
- Often used before a comparison of means

$$W = \frac{(N - k)}{(k - 1)} \cdot \frac{\sum_{i=1}^k N_i (Z_{i\cdot} - Z_{..})^2}{\sum_{i=1}^k \sum_{j=1}^{N_i} (Z_{ij} - Z_{i\cdot})^2},$$

where

- k is the number of different groups to which the sampled cases belong,
- N_i is the number of cases in the i th group,
- N is the total number of cases in all groups,
- Y_{ij} is the value of the measured variable for the j th case from the i th group,
- $Z_{ij} = \begin{cases} |Y_{ij} - \bar{Y}_{i\cdot}|, & \bar{Y}_{i\cdot} \text{ is a mean of the } i\text{-th group,} \\ |Y_{ij} - \tilde{Y}_{i\cdot}|, & \tilde{Y}_{i\cdot} \text{ is a median of the } i\text{-th group.} \end{cases}$

Sample size to conduct hypothesis testing

- One sample

$$n = \left(\frac{Z_{1-\alpha/2} + Z_{1-\beta}}{ES} \right)^2$$
$$ES = \frac{|\mu_1 - \mu_2|}{\sigma}$$

- One sample proportion

$$n = \left(\frac{Z_{1-\alpha/2} + Z_{1-\beta}}{ES} \right)^2$$
$$ES = \frac{p_1 - p_0}{\sqrt{p_1(1-p_1)}}$$

- Two independent sample

$$n_i = 2 \left(\frac{Z_{1-\alpha/2} + Z_{1-\beta}}{ES} \right)^2$$
$$ES = \frac{|\mu_1 - \mu_2|}{\sigma}$$

- Two ind. Sample - proportion

$$n_i = 2 \left(\frac{Z_{1-\alpha/2} + Z_{1-\beta}}{ES} \right)^2$$
$$ES = \frac{|p_1 - p_2|}{\sqrt{p(1-p)}}$$

- Two paired sample


$$n = \left(\frac{Z_{1-\alpha/2} + Z_{1-\beta}}{ES} \right)^2$$
$$ES = \frac{\mu_d}{\sigma_d}$$

ANOVA - Analysis of Variance

- ANOVA is a statistical technique that is used to check if the means of two or more groups are significantly different from each other.
- When we have only two samples, t-test and ANOVA give the same results
- If we conduct multiple t-tests for comparing more than two samples, it will have a compounded effect on the error rate of the result
- ANOVA uses F-tests to statistically test the equality of means.

$$H_0: \mu_0 = \mu_1 = \dots = \mu_m$$

H_1 : at least one μ is unequal



Hypothesis testing: **example**

Problem: We measure the heights of 40 randomly chosen men, and get a mean height of 175cm. We have already known that standard deviation is 20cm. With alpha 5%, can we conclude that mean height of ALL men is greater than 170cm?

$$H_0: \mu = 170$$

$$H_1: \mu > 170$$


Critical value: $\mathbf{z} > 1.645$

Test statistic: $(175-170)/(20/\sqrt{40}) = 1.581$

Test statistic is less than critical value, so we failed to reject H_0 , and conclude that there are not sufficient evidences that population mean of men's height is more than 170cm



Correlation






Correlation

- Correlation is a statistical technique that can show whether and how strongly pairs of numerical variables are related
- Denoted by a **correlation coefficient** (or "r"). It ranges from -1.0 to +1.0.
- The closer r is to +1 or -1, the more closely the two variables are related

Pearson's correlation:

$$r = \frac{\sum (x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum (x_i - \bar{x})^2 \sum (y_i - \bar{y})^2}}$$




Association

- The **Chi-Square Test** of Independence determines whether there is an association between categorical variables
- This test utilizes a contingency table (also known as *cross-tabulation*, *crosstab*, or *two-way table*) to analyze the data

	Smartphones	Laptops	Total Sales (in \$M)
Region 2	132	14	146
Region 3	92	16	108
Total Sales (in \$M)	224	30	254


Correlation does not imply causation



- Seeing two variables moving together does not necessarily mean we know whether one variable causes the other to occur.
- It may be the result of random chance, where the variables appear to be related, but there is no true underlying relationship
- There may be a third, lurking variable that makes the relationship appear stronger (or weaker) than it actually is
- ...but with well-designed empirical research, we can establish causation!
- **So, how do we explore causation? With the right kind of investigation!**



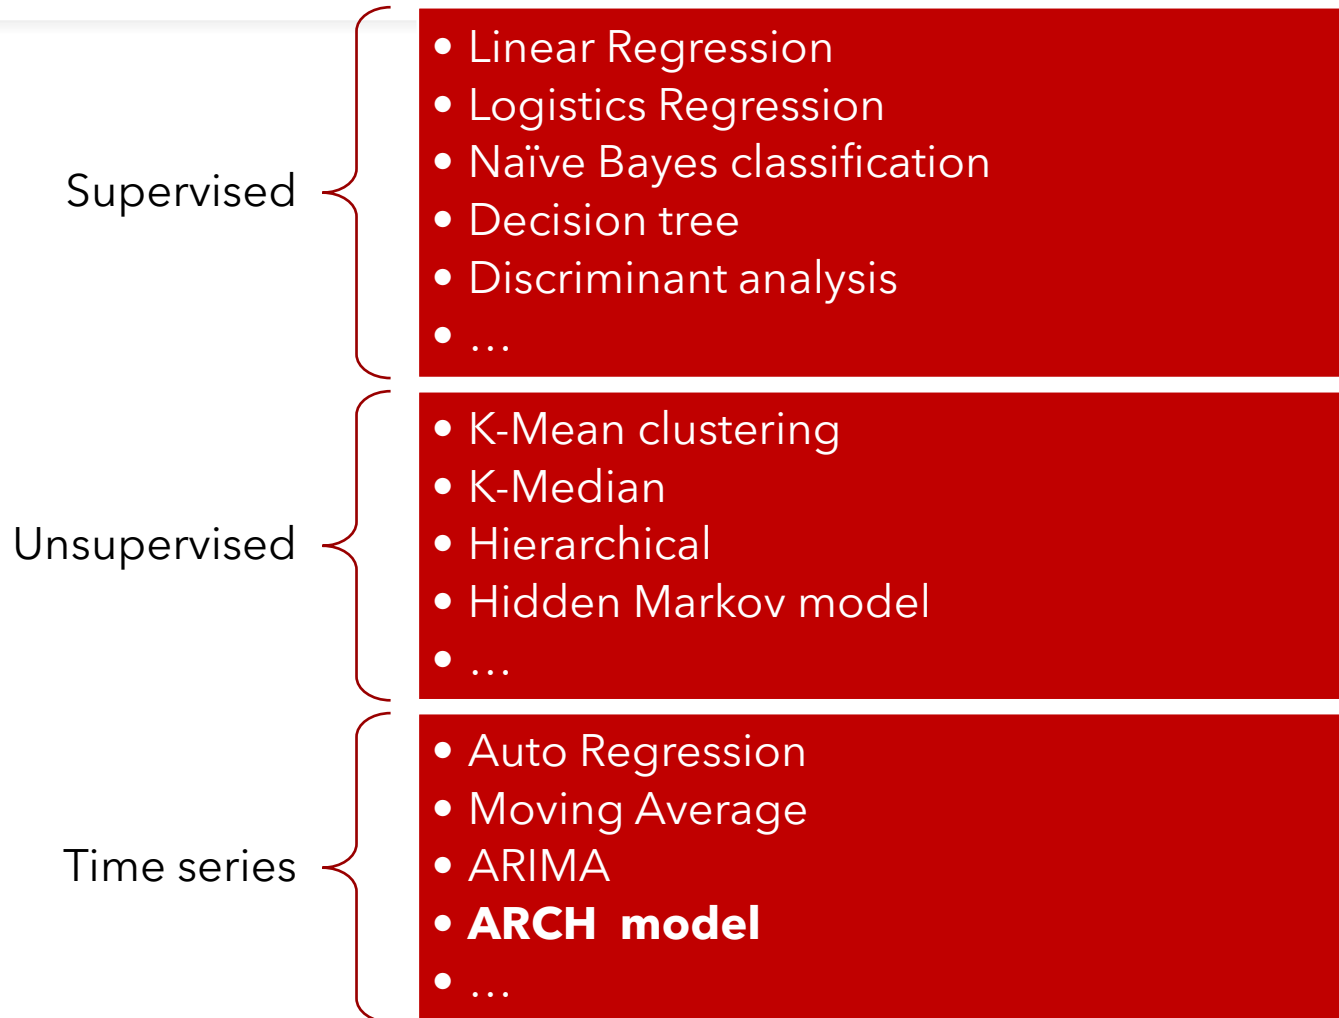
Regression




Statistical modeling

The majority of the problems in statistical inference can be considered to be problems related to statistical modeling

-- Konishi & Kitagawa (2008)




Regression



- Linear regression is a basic and commonly used type of predictive analysis.
- In regression, we examine:
 - does a set of predictor variables do a good job in predicting an outcome (dependent) variable?
 - Which variables in particular are significant predictors of the outcome variable, and in what way do they impact the outcome variable?

Regression

- Regression is one types of statistical modeling
- The **goal** of is to adjust the values of the model's parameters to find the line or curve that comes closest to your data



Simple:

$$y = b_0 + b_1 x + \textit{error}$$

Multiple:

$$y = b_0 + b_1 x_1 + b_2 x_2 + \dots + \textit{error}$$

- We then use statistical inference to test whether regression coefficients (b_0 , b_1 , b_2 , ...) are significant

Regression

- Regression only accepts numerical variables for both **predictors** and **response**
- Categorical variable can be use as predictors after be transformed into **dummy variables**
- Significance of whole model → **F** test (ANOVA), Significance of each predictors → **t**-test
- **R-Squared (or Adjusted R-Squared)** is commonly used to quantify the goodness of the model

Regression

Assumptions for linear regression and ANOVA

- Each group sample is drawn from a normally distributed population
- All populations have a common variance (homoscedasticity)
- All samples are drawn independently of each other
- Within each sample, the observations are sampled randomly and independently of each other (no auto-correlation)
- No correlation between predictors (no multicollinearity)

Statistics vs. Machine Learning

- The two are highly related and share some underlying machinery, but they are different:

	Statistics	Machine Learning
Focus	<u>Building models</u>	<u>Creating system</u> that learn from data
Purpose	<u>Inference, relationship</u> between variables	<u>Prediction</u> accuracy, optimization
Prior assumption about data	Some knowledge (<u>assumption</u>) about population usually required	<u>Without assumption</u>
Dimensionality of data	Usually applied to low-dimensional data	Usually applied to high-dimensional data
Knowledge overlap	No machine learning knowledge required	Some stats knowledge usually needed. Stats is basis for algorithm



Q/A

